



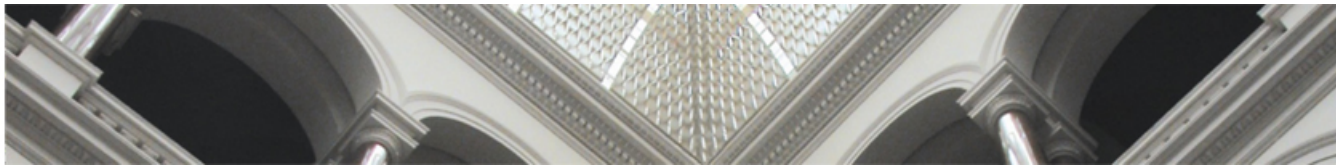
Dimensionality Reduction

Sebastian Schelter | Database Group, TU Berlin | Scalable Data Science: Systems and Methods



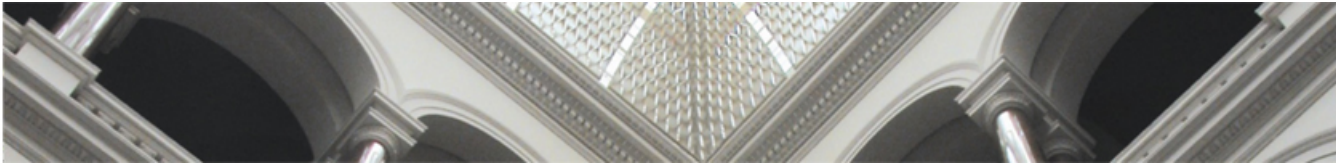
Overview

- Representing data as matrices
- Intuition
- Applications
 - Latent Semantic Indexing
 - Estimating the number of triangles in a network
- Singular Value Decomposition at Scale
- Summary



Overview

- **Representing data as matrices**
- Intuition
- Applications
 - Latent Semantic Indexing
 - Estimating the number of triangles in a network
- Singular Value Decomposition at Scale
- Summary



Representing data as matrices

How would you represent the following types of data as a matrix ?

What would be the rows and columns of the matrix?

What would be the contents of a cell?

- **a collection of text documents**
- **a large network**
- **a set of movie ratings**



Representing data as matrices

	idea	rows	columns	cells
text documents	'bag of words'	documents	terms	term occurrence, term frequency, tf-idf scores
network	adjacency matrix	vertices	vertices	edge occurrences, edge weights
movie ratings	interaction matrix	users	items	occurrence of interaction, rating

- resulting matrices are **high-dimensional** and **extremely sparse** in many cases (problematic!)
- **dimensionality reduction**: reduce data to the “**interesting**” dimensions



Overview

- Representing data as matrices
- **Intuition**
- Applications
 - Latent Semantic Indexing
 - Estimating the number of triangles in a network
- Singular Value Decomposition at Scale
- Summary




Intuition

$$\begin{bmatrix} 2 & 0 \\ 1 & 3 \end{bmatrix}$$



Intuition

$$\begin{bmatrix} 2 & 0 \\ 1 & 3 \end{bmatrix}$$






Intuition

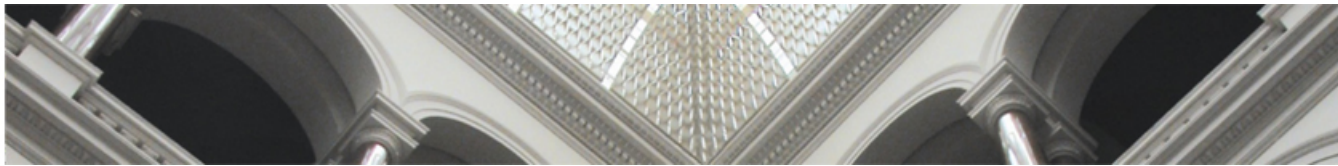
$$\begin{bmatrix} 0.5 & 1 \\ 1.5 & 3 \end{bmatrix}$$



Intuition

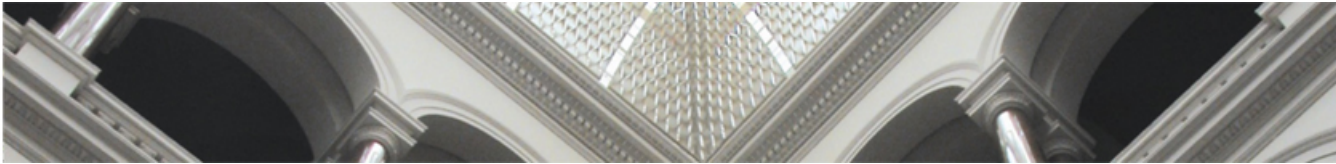
$$\begin{bmatrix} 0.5 & 1 \\ 1.5 & 3 \end{bmatrix}$$






Intuition

$$\begin{bmatrix} 0.6 & 1 \\ 1.5 & 2.9 \end{bmatrix}$$



Intuition

$$\begin{bmatrix} 0.4 & 1 \\ 1.5 & 2.9 \end{bmatrix}$$





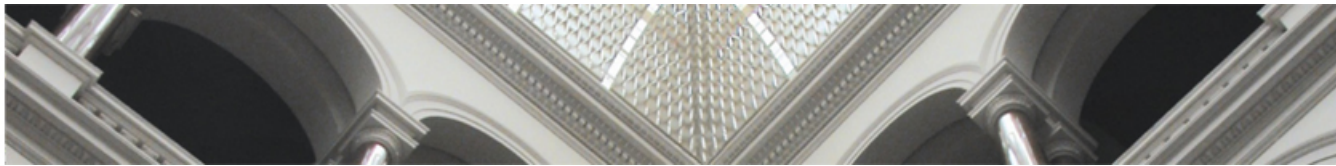
Overview

- Representing data as matrices
- Intuition
- Applications
 - **Latent Semantic Indexing**
 - Estimating the number of triangles in a network
- Singular Value Decomposition at Scale
- Summary



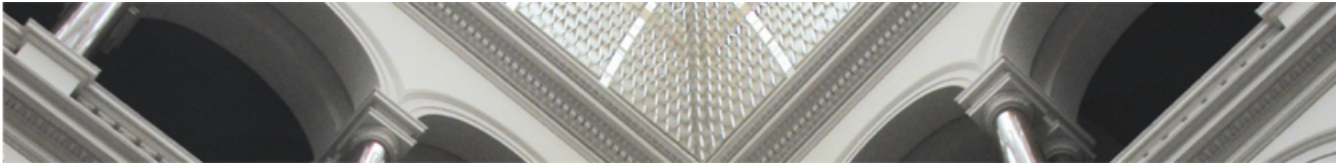
Latent Semantic Indexing

- quick **review of search engines**:
 - documents represented as collection of terms
 - search engines operate on “**inverted index**” from terms to documents
 - search: lookup documents for terms contained in query from inverted index
- mathematically: documents and queries represent vectors in high-dimensional term space (“**vector space model**”)
 - searching means finding the closest document vectors to the query vector
- main **drawback**
 - relies on **lexical matches**, unable to identify synonyms and conceptually close terms



Lexical matching

- imagine a **corpus** with the following three documents
 - document 1: *“bike”*
 - document 2: *“bike harley”*
 - document 3: *“berlin”*
- a query for *“harley”* in a search engine which uses lexical matching only returns document 2
 - yet, document 1 might be relevant as well!
- can we build a search engine that is “smart” enough to also return document 1?



Manual solution: query expansion

- create custom **taxonomies of weighted relations between terms**
e.g. „harley → bike 0.5“
- **automatically expand queries**
 - query “harley” becomes „harley bike^{0.5}“
- **drawbacks**
 - crafting these lists is a lot of work, as they are domain-dependent!
 - might lead to very long queries (expensive!)
 - result quality is hard to predict



Towards Latent Semantic Indexing

- intuition: **structure contained in the corpus** describing relations between terms and documents
- assume **terms and documents** belong to “**latent**” **concepts**, then:
 - a single term describing a particular concept will occur in documents about that concept
 - terms describing the same concept co-occur in documents about that concept
 - documents about a particular concept share a set of characteristic terms

The Linear Algebra view of search

- simplified model:
 - **corpus is represented as document \times term matrix**
 - a cell m,n is 1 if document m contains term n and 0 otherwise

		<i>bike</i>	<i>harley</i>	<i>berlin</i>
	<i>doc1</i>	1	0	0
$A =$	<i>doc2</i>	1	1	0
	<i>doc3</i>	0	0	1

- **queries „harley“ and „harley bike“ are just vectors in the term space** (analogous to documents)

	<i>bike</i>	<i>harley</i>	<i>berlin</i>		<i>bike</i>	<i>harley</i>	<i>berlin</i>
$q_1 =$	0	1	0	$q_2 =$	1	1	0



Search as matrix-vector multiplication

- use the **number of shared terms as similarity measure** between queries and documents
→ **search becomes matrix-vector multiplication**

$$A q^T$$

- examples: search for „harley“ and „harley bike“ in corpus

doc1 : „bike“

doc2: „harley bike“

doc3: „berlin“

Search as matrix-vector multiplication

- example: search for *“harley”*

		bike	harley	berlin
A =	doc1	1	0	0
	doc2	1	1	0
	doc3	0	0	1

$$q_1 = \begin{matrix} & \textit{bike} & \textit{harley} & \textit{berlin} \\ 0 & 1 & 0 \end{matrix}$$

$$A q_1^T = \begin{bmatrix} 1 & 0 & 0 \\ 1 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}$$

Search as matrix-vector multiplication

- example: search for “*harley*”

		bike	harley	berlin
A =	doc1	1	0	0
	doc2	1	1	0
	doc3	0	0	1

$$q_1 = \begin{matrix} & \textit{bike} & \textit{harley} & \textit{berlin} \\ & 0 & 1 & 0 \end{matrix}$$

$$A q_1^T = \begin{bmatrix} 1 & 0 & 0 \\ 1 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix} = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}$$

Search as matrix-vector multiplication

- example: search for “*harley*”

		bike	harley	berlin
A =	doc1	1	0	0
	doc2	1	1	0
	doc3	0	0	1

$$q_1 = \begin{matrix} & \textit{bike} & \textit{harley} & \textit{berlin} \\ & 0 & 1 & 0 \end{matrix}$$

$$A q_1^T = \begin{bmatrix} 1 & 0 & 0 \\ 1 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix} = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}$$

lexical matches in doc1

lexical matches in doc2

lexical matches in doc3

Search as matrix-vector multiplication

- example: search for “*bike harley*”

		bike	harley	berlin
A =	doc1	1	0	0
	doc2	1	1	0
	doc3	0	0	1

$$q_2 = \begin{matrix} & \textit{bike} & \textit{harley} & \textit{berlin} \\ & 1 & 1 & 0 \end{matrix}$$

$$A q_2^T = \begin{bmatrix} 1 & 0 & 0 \\ 1 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix}$$

Search as matrix-vector multiplication

- example: search for “*bike harley*”

		bike	harley	berlin
A =	doc1	1	0	0
	doc2	1	1	0
	doc3	0	0	1

$$q_2 = \begin{matrix} & \textit{bike} & \textit{harley} & \textit{berlin} \\ & 1 & 1 & 0 \end{matrix}$$

$$A q_2^T = \begin{bmatrix} 1 & 0 & 0 \\ 1 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 0 \end{bmatrix}$$

Search as matrix-vector multiplication

- example: search for “*bike harley*”

		bike	harley	berlin
A =	doc1	1	0	0
	doc2	1	1	0
	doc3	0	0	1

$$q_2 = \begin{matrix} & \textit{bike} & \textit{harley} & \textit{berlin} \\ & \mathbf{1} & \mathbf{1} & \mathbf{0} \end{matrix}$$

$$A q_2^T = \begin{bmatrix} \mathbf{1} & \mathbf{0} & \mathbf{0} \\ \mathbf{1} & \mathbf{1} & \mathbf{0} \\ \mathbf{0} & \mathbf{0} & \mathbf{1} \end{bmatrix} \begin{bmatrix} \mathbf{1} \\ \mathbf{1} \\ \mathbf{0} \end{bmatrix} = \begin{bmatrix} \mathbf{1} \\ \mathbf{2} \\ \mathbf{0} \end{bmatrix}$$

lexical matches in doc1

lexical matches in doc2

lexical matches in doc3



Exploring our corpus with Linear Algebra

- vector space model of information retrieval
 - a query is just a vector in term space, analogous to a document
 - we can compute the similarity of this vector to all documents
- **how does this help with finding „latent concepts“ ?**




Exploring our corpus with Linear Algebra

- vector space model of information retrieval
 - a query is just a vector in term space, analogous to a document
 - we can compute the similarity of this vector to all documents
- **how does this help with finding „latent concepts“ ?**
 - **we can compute similarities between documents!**

Exploring our corpus with Linear Algebra

- computing AA^T gives a matrix of document similarities
- cell $A_{m,n}$ holds the number of terms shared between documents m and n
 - documents 1 and 2 are similar!


$$AA^T = \begin{array}{cc} & \begin{array}{ccc} doc1 & doc2 & doc3 \end{array} \\ \begin{array}{c} doc1 \\ doc2 \\ doc3 \end{array} & \begin{array}{ccc} 1 & 1 & 0 \\ 1 & 2 & 0 \\ 0 & 0 & 1 \end{array} \end{array}$$



Exploring our corpus with Linear Algebra

- computing $A^T A$ gives a matrix of term co-occurrences
- cell $A_{m,n}$ holds the number of documents in which terms m and n occur together
→ “harley” and “bike” are related!

		<i>bike</i>	<i>harley</i>	<i>berlin</i>
$A^T A =$	<i>bike</i>	2	1	0
	<i>harley</i>	1	1	0
	<i>berlin</i>	0	0	1




Singular Value Decomposition (SVD)

- **Singular Value Decomposition** of a real $m \times n$ matrix A :
 - U ($m \times m$) and V ($n \times n$) are orthogonal, Σ ($m \times n$) is diagonal
 - Σ has the square roots of the eigenvalues of $A^T A$ and $A A^T$ on its diagonal in descending order (**singular values**)
 - columns of U are the corresponding eigenvectors of $A A^T$ (**left singular vectors**)
 - columns of V are the corresponding eigenvectors of $A^T A$ (**right singular vectors**)
 - if we only keep the top k singular values of A , we get the optimal rank k approximation A_k of A

$$A = U \Sigma V^T$$

$$A_k = U_k \Sigma_k V_k^T$$



Interpreting the SVD

- examine the **rank-2 decomposition of A**
 - rows of A (documents) and columns of A (terms) projected onto a 2-dimensional space, the **latent concept space**
 - „bike“ and „harley“ as well as *doc1* and *doc2* point into the same direction („berlin“ and *doc3* point into perpendicular directions)

$$U_2 = \begin{matrix} doc1 \\ doc2 \\ doc3 \end{matrix} \begin{bmatrix} -.53 & 0 \\ -.85 & 0 \\ 0 & 1 \end{bmatrix} \quad \Sigma_2 = \begin{bmatrix} 1.62 & 0 \\ 0 & 1 \end{bmatrix} \quad V_2 = \begin{matrix} bike \\ harley \\ berlin \end{matrix} \begin{bmatrix} -.85 & 0 \\ -.53 & 0 \\ 0 & 1 \end{bmatrix}$$

- dimensions of the space loosely correspond to concepts („motorcycles“ and „berlin“)
- replacement documents and terms with vectors that represent their association to the concepts
- singular values denote the “importance” of the concepts

Interpreting the SVD

- **latent concept space**
 - dimensions represent „concepts“ (might be hard to interpret)
 - conceptually similar documents and terms are near to each other





Search in the concept space

- **search in the concept space**
 - project the query into the concept space (**fold-in**)

$$q = \begin{matrix} & \textit{bike} & \textit{harley} & \textit{berlin} \\ \begin{matrix} \textit{bike} & \textit{harley} & \textit{berlin} \end{matrix} & \begin{bmatrix} 0 & 1 & 0 \end{bmatrix} \end{matrix} \quad \hat{q} = q V \Sigma^{-1} = \begin{bmatrix} -.85 & 0 \end{bmatrix}$$

Search in the concept space (the punchline)

- **search in the concept space**
 - project the query into the concept space (**fold-in**)

$$\begin{array}{c}
 \text{bike} \quad \text{harley} \quad \text{berlin} \\
 q = \begin{bmatrix} 0 & 1 & 0 \end{bmatrix} \quad \hat{q} = q V \Sigma^{-1} = \begin{bmatrix} -.85 & 0 \end{bmatrix}
 \end{array}$$

- compare the projected query to the document concept vectors

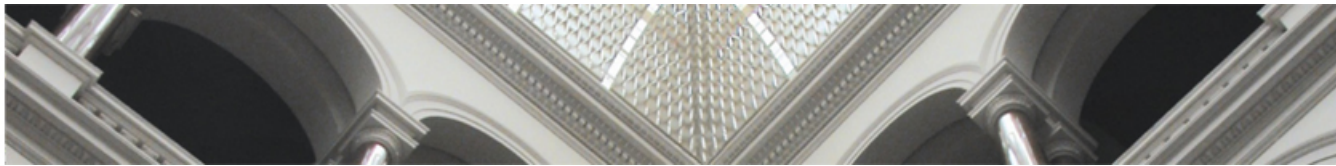
$$U_2 \hat{q}^T = \begin{bmatrix} -.53 & 0 \\ -.85 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} -.85 \\ 0 \end{bmatrix} = \begin{bmatrix} 0.45 \\ 0.72 \\ 0 \end{bmatrix} \begin{array}{l} \text{doc1} \\ \text{doc2} \\ \text{doc3} \end{array}$$

- **query matches document 1 although it does not contain the search term “harley” !!!**



Drawbacks of Latent Semantic Indexing

- computing the SVD of a large corpus is computationally expensive
 - needs **constant re-computation** for new documents
- **hard to scale:**
 - document concept matrix typically dense
 - every document needs to be inspected at query time!
- works well for synonyms but does not handle polysemy




Overview

- Representing data as matrices
- Intuition
- Applications
 - Latent Semantic Indexing
 - **Estimating the number of triangles in a network**
- Singular Value Decomposition at Scale
- Summary


Social Graphs and Triangles

- **social graphs**
 - vertices represent users
 - edges represent connections between users, e.g. friendship, following, etc

- a **triangle** in a graph is a triple of interconnected vertices (*3-clique*)




- social graphs grow by “**closing triangles**”:
 - “becoming friends with a friend of a friend”




Local clustering coefficient

- a measure of the local “**connectedness**” of a vertex
 - the number of links between the neighbors of a vertex i divided by the maximum possible number of links between these neighbors

$$C_i = \frac{2|\{e_{jk} : v_j, v_k \in N_i, e_{jk} \in E\}|}{k_i(k_i - 1)}$$



local clustering coefficient of 1




local clustering coefficient of 1/3


Local clustering coefficient

- local clustering coefficient C_i of a vertex i is connected to the **number of triangles** t_i which vertex i is a part of

$$C_i = \frac{2|\{e_{jk} : v_j, v_k \in N_i, e_{jk} \in E\}|}{k_i(k_i - 1)} = \frac{2t_i}{k_i(k_i - 1)}$$



local clustering coefficient of 1



local clustering coefficient of 1/3




Applications of the local clustering coefficient

- statistic of interest in **network science**
- can be used to **identify spammers amongst high degree vertices in a social network**
 - spammers should have a local clustering coefficient below average
as they typically randomly connect to other users

Adjacency matrix of a graph


- A_{ij} is 1 if there is an edge between vertices i and j , 0 otherwise



$$A = \begin{bmatrix} 0 & 1 & 1 & 1 & 1 \\ 1 & 0 & 1 & 1 & 0 \\ 1 & 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 & 0 \end{bmatrix}$$

Powers of the adjacency matrix


- **powers of the adjacency matrix give information about paths in a network**
- e.g., cell (i,j) of A^2 holds the number of paths of length 2 between vertex i and vertex j



$$A^2 = \begin{bmatrix} 4 & 2 & 1 & 2 & 1 \\ 2 & 3 & 1 & 1 & 2 \\ 1 & 1 & 2 & 2 & 1 \\ 2 & 1 & 2 & 3 & 1 \\ 1 & 2 & 1 & 1 & 2 \end{bmatrix}$$


Powers of the adjacency matrix

- a **triangle is a path of length 3** from a vertex back to itself
→ the diagonal of A^3 holds the number of triangles* for each vertex
- unfortunately, multiplying large matrices quickly becomes infeasible



$$A^3 = \begin{bmatrix} 6 & & & & \\ & 4 & & & \\ & & 2 & & \\ & & & 4 & \\ & & & & 2 \end{bmatrix}$$

*(actually, its twice the number of triangles, as we count each triangle twice)



Diagonalization

- **Diagonalization** of a diagonalizable square $n \times n$ matrix A :
 - $Q (n \times n)$ is orthogonal, $\Delta (n \times n)$ is diagonal
 - Δ has the **eigenvalues** of A on its diagonal in descending order
 - Q has the corresponding eigenvectors of A as columns
 - if we only keep the top k eigenvalues of A , we get the optimal rank k approximation A_k of A

$$A = Q \Delta Q^T$$



Diagonalization and the powers of a matrix

- diagonalization provides an easy way to compute the powers of a matrix


$$A = Q \Delta Q^T$$

- after diagonalization of A, **powers of A can be computed by computing only the powers of Δ** (which is easy since Δ is a diagonal matrix)

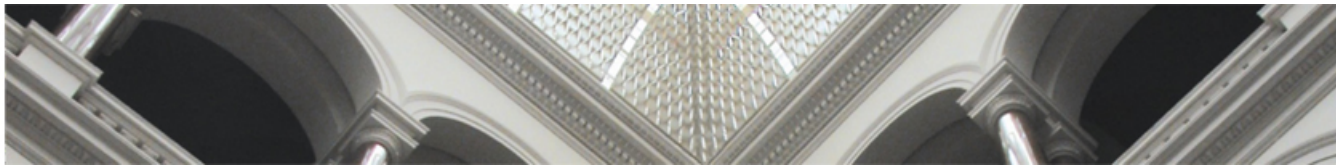
$$A^3 = Q \Delta Q^T Q \Delta Q^T Q \Delta Q^T = Q \Delta^3 Q^T$$

Powers of the adjacency matrix

- using only a few eigenvectors and eigenvalues usually suffices to get a good estimate of the number of triangles



$$A^3 \approx Q_3 \Delta_3^3 Q_3^T = \begin{bmatrix} 6.0093 & & & & \\ & 3.9931 & & & \\ & & 1.9937 & & \\ & & & 3.9931 & \\ & & & & 1.99 \end{bmatrix}$$




Overview

- Representing data as matrices
- Intuition
- Applications
 - Latent Semantic Indexing
 - Estimating the number of triangles in a network
- **Singular Value Decomposition at Scale**
- Summary



Singular Value Decomposition

- if matrix A only has a small number of columns:
 - compute $A^T A$ at scale
 - use a fast single machine solver to compute the first k eigenvalues and eigenvectors of $A^T A$
- this gives Σ_k and V_k
- compute U_k at scale as $A V_k \Sigma_k^{-1}$



Lanczos algorithm

- Lanczos algorithm
 - computes eigendecomposition of matrix A
 - SVD can be computed from eigendecomposition of $A^T A$
- Krylov subspace method, repeatedly multiplies matrix with a random vector
- creates a tridiagonal matrix T which has the same eigenvalues as A (and makes them easy to find)
- easy to scale as it requires only matrix-vector multiplications

```

 $v_1 \leftarrow$  random vector with norm 1
 $v_0 \leftarrow 0$ 
 $\beta_1 \leftarrow 0$ 

```

```

for  $j = 1, 2, \dots, m-1$ 
   $w_j \leftarrow A v_j$ 
   $\alpha_j \leftarrow w_j^T v_j$ 
   $w_j \leftarrow w_j - \alpha_j v_j - \beta_j v_{j-1}$ 
   $\beta_{j+1} \leftarrow |w_j|$ 
   $v_{j+1} \leftarrow w_j / \beta_{j+1}$ 

```

```

 $w_m \leftarrow A v_m$ 
 $\alpha_m \leftarrow w_m^T v_m$ 

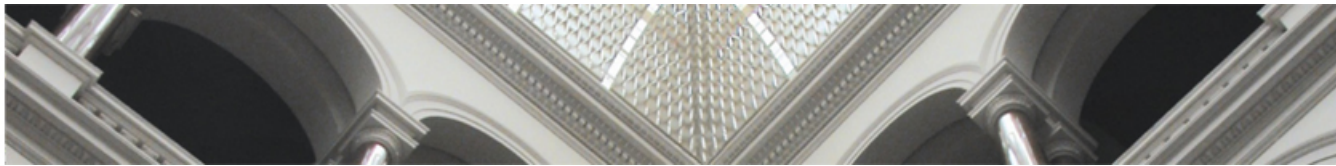
```

$$T = \begin{pmatrix} \alpha_1 & \beta_2 & 0 & 0 & 0 & 0 \\ \beta_2 & \alpha_2 & \beta_3 & 0 & 0 & 0 \\ 0 & \beta_3 & \alpha_3 & \dots & 0 & 0 \\ 0 & 0 & \dots & \dots & \beta_{m-1} & 0 \\ 0 & 0 & 0 & \beta_{m-1} & \alpha_{m-1} & \beta_m \\ 0 & 0 & 0 & 0 & \beta_m & \alpha_m \end{pmatrix}$$



Overview

- Representing data as matrices
- Intuition
- Applications
 - Latent Semantic Indexing
 - Estimating the number of triangles in a network
- Singular Value Decomposition at Scale
- **Summary**



Summary

- often: **data looks high dimensional, but has low-rank structure**
- **dimensionality reduction projects data onto a low-rank space**
 - via matrix decomposition techniques, e.g. SVD, Eigendecomposition
- many interesting problems solvable via generalization provided by dimensionality reduction
 - Latent Semantic Indexing
 - Triangle Count Estimation
 - Latent Factor Models for Recommender Systems (coming up)
- SVD at Scale
 - use **Lanczos algorithm** (main operation: matrix-vector multiplications)
- not covered:
 - **Principal Components Analysis (PCA)**, as it is hard to scale



Further Reading

- Manning, C. D., Raghavan, P., & Schütze, H. (2008). *Introduction to information retrieval* (Vol. 1, p. 496). Cambridge: Cambridge university press.
- Dumais, S. T. (2004). *Latent semantic analysis*. Annual review of information science and technology, 38(1), 188-230.
- Kang, U., Meeder, B., & Faloutsos, C. (2011). *Spectral analysis for billion-scale graphs: Discoveries and implementation*. In Advances in Knowledge Discovery and Data Mining (pp. 13-5). Springer Berlin Heidelberg.
- Strang, G. (1993). *The fundamental theorem of linear algebra*. American Mathematical Monthly, 848-855.