

SystemML: Declarative Machine Learning

By Ahmed Tidjani and Lorenzo Toso



Fachgebiet Datenbanksysteme und Informationsmanagement
Technische Universität Berlin

<http://www.dima.tu-berlin.de/>

- What is SystemML
- History
- Design Principles
- Technical Details
- Usage
- Comparison to similar frameworks
- Conclusion
- References

■ Regular workflow:

- Data scientist develops algorithm in R or Python
- Tests it on small dataset
- Systems engineer ports algorithm to Scala
 - Big Data operations are possible
- Possibly multiple iterations required!



[1] systemml.apache.org

■ SystemML:

- Defines a Declarative High-Level Language for machine learning (DML)
 - Similar Syntax to Python or R
- System running on Hadoop, Spark, Standalone or technically any other BigData-Framework
 - Flink integration is being worked on (First tests by Andreas Kunft)
- Allows automatic scaling of machine learning algorithms



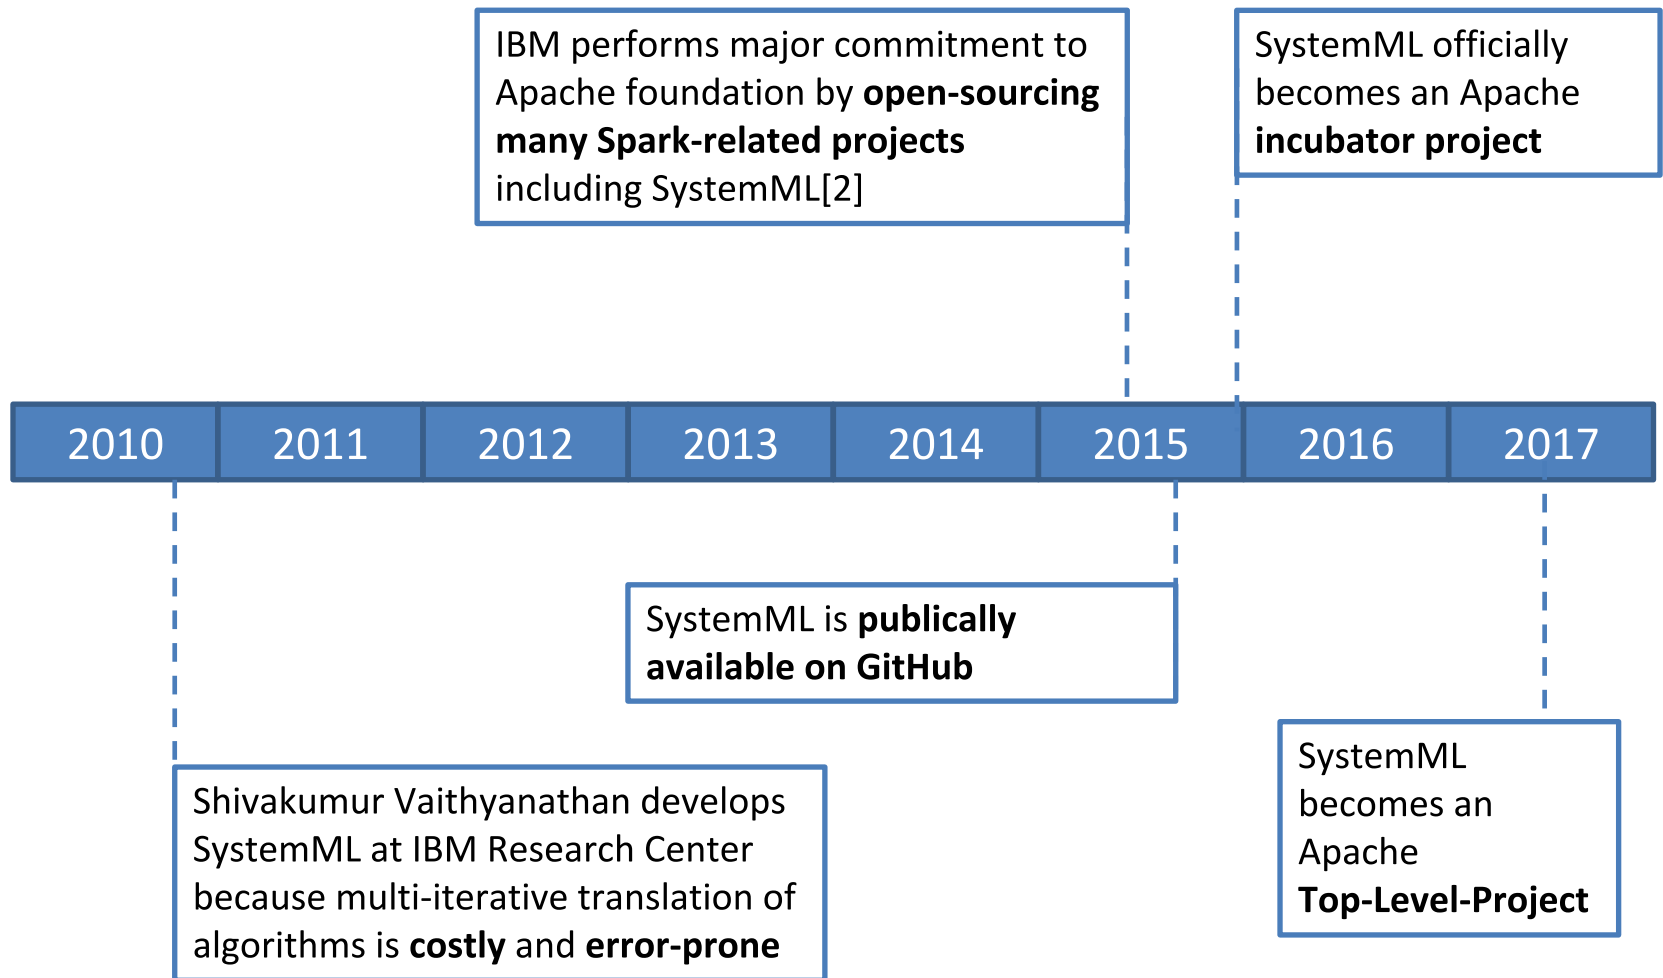
[1] systemml.apache.org

■ SystemML workflow:

- Data scientists declare algorithms like they are used to
 - Test them on a single machine with small datasets
 - Test them on a cluster with large datasets
 - Deploy them for real applications
- ⇒ Eliminates need error-prone translation of algorithms



[1] systemml.apache.org



- Declarative Machine Learning (DML)
- Syntax similar to R (DML) or Python (PyDML)
- Translation from Python+Numpy is straightforward
- Contains a rich set of functions:
 - matrix operations
 - statistical methods
 - machine learning algorithms
- Control sequences like „parfor“-loops for task-parallel operations

■ DML code example

```

script = """
    # add constant feature to X to model intercepts
    X = cbind(X, matrix(1, rows=nrow(X), cols=1))
    max_iter = 100
    w = matrix(0, rows=ncol(X), cols=1)
    for(i in 1:max_iter){
        XtX = t(X) %*% X
        dw = XtX %*% w - t(X) %*% y
        alpha = -(t(dw) %*% dw) / (t(dw) %*% XtX %*% dw)
        w = w + dw*alpha
    }
    bias = as.scalar(w[nrow(w),1])
    w = w[1:nrow(w)-1,]
"""
    
```


- SystemML aims to hide all the technical details from the user
- Transparency regarding the underlying framework
 - Spark, MapReduce, Standalone
 - Optimizations for individual frameworks
 - E.g. Spark specific optimization of caching algorithms
- Optimizes execution parameters based on data and cluster characteristics
 - E.g. decides whether to use a cluster at all or if things should be done on a single machine
 - Tries to minimize shuffling of data

- All input data is converted into binary block matrices
 - Other systems use a similar approach [5]
- Map pair of indices to Matrix-Content
 - (long|long) -> Matrix-Content
 - Allows arbitrary distribution
- All blocks are squared and of fixed size
 - default: 1000x1000

⇒ 8MB fits in regular L3 Cache
- Blocks decide independently whether they are
 - Empty (Simple flag is stored)
 - Dense (All values are stored)
 - Sparse (Store either row or column index - Value pairs)
 - Ultra-Sparse (Store row-column-value pairs)

- Caching is completely hidden from the user
 - Needs to be done well automatically

- Uses storage level MEMORY_AND_DISK
 - RDDs are stored as deserialized Objects in the JVM on Memory/Disk
 - Except for „Ultra-Sparse-Matrices“. These use MEMORY_AND_DISK_SER

- Uses Spark Cache and Checkpoints
 - Cache: Materializes RDDs and keeps them in memory
 - Including Lineage for fault tolerance
 - Checkpoint: Saves RDDs to HDFS
 - without lineage
 - HDFS is fault tolerant by replication

- Performs fixed set of caching optimizations
 - Checkpoint injection (brings matrices into read-optimized form)
 - Before loops
 - After persistent reads
 - Deferred optimization of parfor-blocks
 - Removing of redundant checkpoints

- Shuffling of data is very expensive
 - Needs to be avoided
- Use of operations that avoid shuffling is highly preferred
 - “Partition-Preserving operations” are used
 - Key must not change
 - More restrictive use of API
- Introduce explicit repartitioning to avoid unnecessary reshuffling
- Use Partition-Exploiting operations
 - E.g. Transposing on a single node

- Basic Spark execution type
 - Simplest approach for choosing execution type
 - Worst case memory estimate for an operation is greater than the available memory on a single node
 - ⇒ Run operation on spark

- Transitive Spark execution type
 - Data would fit in memory of a single machine
 - Data transfer is reduced by scheduling Spark-Job
 - Example:
 - Sum over many values
 - Every node can compute the sum of the values it already has in memory
 - Only intermediate results need to be transferred

- Optimization of High-Level-Operators (HOPs)
 - Algebraic rewrites
 - Operator order
 - E.g. Matrix multiplication of sparse matrices
 - Compute memory estimates

- Compute Low-Level-Operators (LOPs)
 - Based on memory estimates
 - Based on cluster characteristics
 - Have corresponding runtime implementations

- SystemML performs dynamic recompilation
 - Between major program blocks
 - When the optimizer explicitly injects recompilation points
- Allows to adapt runtime execution plan to previously unknown data characteristics
- Problem Spark Lazy Evaluation:
 - Some necessary values may not be computed yet
 - ⇒ Characteristics of operations are used to determine output sizes and data types when possible

■ From a Scala-Program

```
import org.apache.sysml.api.mlcontext._
import org.apache.sysml.api.mlcontext.ScriptFactory._
val ml = new MLContext(sc)
```

■ From a Spark-Shell

```
scala> import org.apache.sysml.api.mlcontext._
import org.apache.sysml.api.mlcontext._
```

```
scala> import org.apache.sysml.api.mlcontext.ScriptFactory._
import org.apache.sysml.api.mlcontext.ScriptFactory._
```

```
scala> val ml = new MLContext(sc)
```

Welcome to Apache SystemML!

```
ml: org.apache.sysml.api.mlcontext.MLContext = org.apache.sysml.api.mlcontext.MLContext@12139db0
```

- From Python
- Or a Jupyter-Notebook

```
import systemml as sml
import numpy as np
m1 = sml.matrix(np.ones((3,3)) + 2)
m2 = sml.matrix(np.ones((3,3)) + 3)
m2 = m1 * (m2 + m1)
m4 = 1.0 - m2
m4.sum(axis=1).toNumPy()
```

- As Standalone using DML

```
$ ./runStandaloneSystemML.sh scripts/utils/sample.dml -nvars X=data/haberman.data sv=data/perc.csv O=data/haberman.part ofmt="csv"
```

- Download Spark binaries
 - <https://spark.apache.org/downloads.html>
- Download SystemML sources
 - <https://systemml.apache.org/download>
- Build sources using mvn install -DskipTests
- Set SPARK_HOME environment variable
- Set SYSTEMML_HOME environment variable

- Launch PySpark in Jupyter notebook with SystemML
 - PYSPARK_DRIVER_PYTHON=jupyter
 - PYSPARK_DRIVER_PYTHON_OPTS="notebook"
 - \$ pyspark --master local[*] --driver-memory 3G --driver-class-path SystemML.jar --jars SystemML.jar

- PySpark integration in Jupyter notebook
- Linear regression algorithm implemented in SystemML DML
- Hybrid execution: Spark + CP
- Bike-sharing data set

Data Set Characteristics:	Univariate	Number of Instances:	17389	Area:	Social
Attribute Characteristics:	Integer, Real	Number of Attributes:	16	Date Donated	2013-12-20
Associated Tasks:	Regression	Missing Values?	N/A	Number of Web Hits:	143207

<http://archive.ics.uci.edu/ml/datasets/Bike+Sharing+Dataset>

- Initializing SystemMLContext with SparkContext

```
from systemml import MLContext, dml, dmlFromResource  
  
ml = MLContext(sc)  
  
print ("Spark Version:" + sc.version)  
print ("SystemML Version:" + ml.version())  
print ("SystemML Built-Time:" + ml.buildTime())
```

- Executing LinearRegDS.dml

```
scriptName = "LinearRegDS.dml"
dml_script = dmlFromResource(scriptName)

prog = dml_script.input(X=train_data[:, :-3], y=train_data[:, -1].reshape(-1, 1)).input('$icpt', 1.0).output('beta_out')

w = ml.execute(prog).get('beta_out')
w = w.toNumPy()
bias=w[-1]
w = w[:-1, :]
test_y = test_data[:, :-3].dot(w)+bias
```

- Extract from LinerRegDS.dml

```
# BEGIN THE DIRECT SOLVE ALGORITHM (EXTERNAL CALL)

A = t(X) %** X;
b = t(X) %** y;
if (intercept_status == 2) {
    A = t(diag (scale_X) %** A + shift_X %** A [m_ext, ]);
    A =    diag (scale_X) %** A + shift_X %** A [m_ext, ];
    b =    diag (scale_X) %** b + shift_X %** b [m_ext, ];
}
A = A + diag (lambda);

print ("Calling the Direct Solver...");

beta_unscaled = solve (A, b);
```

- Output of LinearRegDS.dml script

```

/bin/bash

/bin/bash

[I 17:08:11.644 NotebookApp] Saving file at /Untitled1-Copy1.ipynb
[I 17:10:11.644 NotebookApp] Saving file at /Untitled1-Copy1.ipynb
BEGIN LINEAR REGRESSION SCRIPT
Reading X and Y...
Calling the Direct Solver...
Computing the statistics...
AVG_TOT_Y,191.1264125534282
STDEV_TOT_Y,182.0341548848353
AVG_RES_Y,3.009216507337137E-13
STDEV_RES_Y,142.3986352096892
DISPERSION,20277.371309582133
R2,0.38849414143527206
ADJUSTED_R2,0.38806415958230267
R2_NOBIAS,0.38849414143527206
ADJUSTED_R2_NOBIAS,0.38806415958230267
Writing the output matrix...
END LINEAR REGRESSION SCRIPT

```


- SystemML stats: `ml = ml.setStatistics(True)`

```
Writing the output matrix...
END LINEAR REGRESSION SCRIPT
SystemML Statistics:
Total elapsed time:          0.181 sec.
Total compilation time:      0.000 sec.
Total execution time:        0.181 sec.
Number of compiled Spark inst: 2.
Number of executed Spark inst: 2.
Cache hits (Mem, WB, FS, HDFS): 26/0/0/0.
Cache writes (WB, FS, HDFS): 9/0/0.
Cache times (ACQr/m, RLS, EXP): 0.000/0.002/0.001/0.000 sec.
HOP DAGs recompiled (PRED, SB): 0/0.
HOP DAGs recompile time:     0.000 sec.
Spark ctx create time (lazy): 0.000 sec.
Spark trans counts (par,bc,col):0/0/0.
Spark trans times (par,bc,col): 0.000/0.000/0.000 secs.
Total JIT compile time:      0.084 sec.
Total JVM GC count:          19.
Total JVM GC time:           0.584 sec.
Heavy hitter instructions (name, time, count):
-- 1)   tsmm      0.006 sec      3
-- 2)   ba+*      0.004 sec      2
-- 3)   append    0.003 sec      9
```

- **Property 1: Independence of Data structures**
 - All data types are given as abstractions with no access to physical data
- **Property 2: Independence of Data Flow Properties**
 - User has no control over partitioning, caching, blocking etc.
- **Property 3: Analysis-Centric Operation Primitives**
 - The framework implements basic functions for statistics and ML
- **Property 4: Known Semantics of Operation Primitives**
 - Knowledge of semantics (associativity, sparse-safeness, etc)

- **Property 5: Implementation-Agnostic Operations**
 - The specification of algorithms is independent of the underlying runtime
- **Property 6: Well-Defined Plan Optimizazion Objective**
 - Objective for execution plan optimization is clearly specified
- **Property 7: Implementation-Agnostic Results:**
 - Results of algorithms are equivalent independently of the used runtime
- **Property 8: Deterministic Results:**
 - Results for multiple runs are equivalent

Table 2: Classification of Existing Systems wrt Declarative ML Algorithms (Type 1).

(P1 Indep. Data Structures, P2 Indep. Data Flow Properties, P3 Analysis-Centric Operations, P4 Known Operations, P5 Impl.-Agnostic Operations, P6 Well-Def. Optim. Objective, P7 Impl.-Agnostic Results, P8 Deterministic Results)

Name	Dist.	Basic Properties								Type	Objective
		P1	P2	P3	P4	P5	P6	P7	P8		
RIOT [43]		✓	✓	✓	✓	✓	✓	✓	✓	1	min runtime
OptiML [35]		✓	✓	✓	✓	✓	✓	✓	✓	1	min runtime
SystemML [7, 22]	✓	✓	✓	✓	✓	✓	✓	✓	✓	1	min runtime s.t. memory constraints
Mahout Samsara [15]	✓			✓	✓		✓	✓	✓	N/A	min runtime
Distributed R [39]	✓			✓			✓	✓	✓	N/A	min runtime
Cumulon [24, 25]	✓	✓	✓	✓	✓	✓		✓	✓	N/A	min costs s.t. runtime constraints
DMac [40]	✓		✓	✓	✓		✓	✓	✓	N/A	min runtime s.t. memory constraints
TensorFlow [1]	✓	✓	✓	✓		✓	✓		✓	N/A	min runtime s.t. resource constraints
SciDB [9, 34]	✓	✓	✓	✓	✓	✓	✓	✓	✓	1	min runtime
SimSQL [11]	✓	✓	✓	✓	✓	✓	✓	✓	✓	1	min runtime
ScalOps [8]	✓		✓				✓	✓	✓	N/A	min runtime
Tupleware [13]	✓		✓				✓	✓	✓	N/A	min runtime
Emma [4]	✓		✓				✓	✓	✓	N/A	min runtime

[6] Declarative Machine Learning - A Classification of Basic Properties and Types

- SystemML performs a great job at hiding technical details from the user
 - User only concentrates on developing the algorithm
 - Algorithms is automatically scaled
- The provided DML is very simple and straightforward if you know R or Python
- The framework is relatively new and currently unproven
 - Small community
 - Only academic uses known
- Can heuristics really perform sufficient parameter optimization?
 - Not well tested
 - No representative benchmarks

- [1] systemml.apache.org
- [2] Ghoting, Amol, et al. "SystemML: Declarative machine learning on MapReduce." *Data Engineering (ICDE), 2011 IEEE 27th International Conference on*. IEEE, 2011.
- [3] Matthias Boehm, Michael W. Dusenberry, Deron Eriksson, Alexandre V. Evfimievski, Faraz Makari Manshadi, Niketan Pansare, Berthold Reinwald, Frederick R. Reiss, Prithviraj Sen, Arvind C. Surve, and Shirish Tatikonda. 2016. SystemML: declarative machine learning on spark. *Proc. VLDB Endow.* 9, 13 (September 2016), 1425-1436. DOI: <http://dx.doi.org/10.14778/3007263.3007279>
- [4] Matthias Boehm, Shirish Tatikonda, Berthold Reinwald, Prithviraj Sen, Yuanyuan Tian, Douglas R. Burdick, and Shivakumar Vaithyanathan. 2014. Hybrid parallelization strategies for large-scale machine learning in SystemML. *Proc. VLDB Endow.* 7, 7 (March 2014), 553-564. DOI=<http://dx.doi.org/10.14778/2732286.2732292>
- [5] Bosagh Zadeh, Reza, et al. "Matrix computations and optimization in apache spark." *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. ACM, 2016.
- [6] Boehm, Matthias, et al. "Declarative Machine Learning-A Classification of Basic Properties and Types." *arXiv preprint arXiv:1605.05826* (2016).
- [7] <https://github.com/apache/systemml>