

Distributed Algorithms 2015/16

Flooding, Broadcast and Echo

Reinhardt Karnapke | Communication and Operating Systems Group

Overview

Flooding


- Distribution of Information (e.g. node-ID) with or without confirmation to all nodes using *all edges*

Echo


- Distribution of information to all nodes using all edges with selective confirmation
- Collecting information
- Construction of a spanning tree

Broadcast

- Distribution of information to all nodes with or without acknowledgement with special topologies




Flooding




Information Distribution with Flooding

- Precondition: Connected topology
- Principle:
Each node tells a *new* rumor to *all other* neighbors



Information Distribution with Flooding

- Every node tells a *new* rumor that it got from one of its neighbors to all other neighbors
- Already known rumors are ignored
- Step by step all nodes will be informed




Flooding-Algorithm

```
I: {NOT informed}
    SEND <info> TO all neighbors
    informed := TRUE;
```

```
R: {A message <info> is received}
    IF NOT informed THEN
        SEND <info> TO all other neighbors;
        informed := TRUE;
    FI
```


Initially, `informed == FALSE` for all processes
Action I is carried out by the initiator spontaneously
Are several competing initiators allowed?



Information Distribution with Flooding

How many messages are sent?

- Let n be the number of nodes and e the number of edges




Information Distribution with Flooding

How many messages are sent?


- Let n be the number of nodes and e the number of edges
- Each node sends over all his incident edges
→ $+2e$ messages
- But not back over its activation edge
→ $-n$ messages
- Exception: initiator (has no activation edge)
→ $+1$ message
- ⇒ Altogether $2e - n + 1$ messages

> How does the initiator know that all nodes were reached? → Termination detection (but how?)



Flooding with Confirmation

- Two message types: Explorers and confirmations
- A process acknowledges an explorer with a confirmation, as soon as it has received a confirmation for all explorers sent by itself due to the receiving of that explorer
 - First received explorer (activation edge):
confirmation after arrival of $\#neighbor - 1$ receipts
→ leafs send confirmation immediately
 - Further explorer: confirmation sent immediately
- Algorithm terminates, if the initiator received a confirmation from every neighbor




Flooding with Confirmation I (wrong!)

```
I: {NOT informed} // Executed by Initiator
  SEND Explorer TO all Neighbors;
  informed := TRUE;
```

```
{Explorer from neighbor N is received}
  IF NOT informed THEN
    SEND Explorer TO all Neighbors except N;
    informed := TRUE;
    A := N;
  FI
```

Initially, `informed == false`
and `Count == 0` for all nodes

```
{Confirmation is received}
  Count := Count + 1;
  IF (NOT Initiator) AND (Count == #Neighbors - 1) THEN
    SEND Confirmation TO Neighbor A;
  FI
  IF Initiator AND (Count == #Neighbors) THEN
    Exit; // Algorithm is terminated.
  FI
```




Flooding with Confirmation II (right)

```
I: {NOT informed} // Executed by Initiator
  SEND Explorer TO all Neighbors;
  informed := TRUE;
```


```
{Explorer from neighbor N is received}
  IF NOT informed THEN
    SEND Explorer TO all Neighbors except N;
    informed := TRUE;
    A := N;
  ELSE
    SEND Confirmation TO N;
  FI
```

Initially, informed == false
and Count == 0 for all nodes

```
{Confirmation is received}
  Count := Count + 1;
  IF (NOT Initiator) AND (Count == #Neighbors - 1) THEN
    SEND Confirmation TO Neighbor A;
  FI
  IF Initiator AND (Count == #Neighbors) THEN
    Exit; // Algorithm is terminated.
  FI
```




Flooding with Confirmation - Example




Here, the number of received confirmations is counted.

Flooding with Confirmation - Example



Flooding with Confirmation - Example




Flooding with Confirmation

How many explorers altogether?

How many confirmations altogether?

Altogether?



Flooding with Confirmation


How many explorers altogether?

- Every node sends an explorer on all edges $\rightarrow +2e$ explorer
- But not on its activation edge $\rightarrow -n$ explorer
- Exception initiator $\rightarrow +1$ explorer
- $2e - n + 1$ explorer


How many confirmations altogether?

- Every node gets a confirmation on every edge $\rightarrow +2e$ messages
- But not on its activation edge $\rightarrow -n$ messages
- Exception initiator $\rightarrow +1$ message
- $2e - n + 1$ confirmations

Altogether: $4e - 2n + 2$ messages, double the number for flooding without confirmation




Echo




Echo-Algorithm

- Initially all nodes are *white*
- The unique initiator becomes *red* and sends red messages (explorers) to all its neighbors




Echo-Algorithm

- A white node, receiving an explorer, becomes red itself and memorizes that „first“ edge (*activation edge*) and sends explorers to all its neighbors
- On an edge, where two explorers meet, the cycle is broken (i.e., the explorers are swallowed)




Echo-Algorithm

- A red node which has received an explorer *or* echo over *all* its edges becomes green and sends a green echo over its „first“ edge which also becomes green
- Leafs immediately send an echo when receiving an explorer



Echo-Algorithm

- By and by all nodes and a part of the edges turn green
- The algorithm terminates when the initiator turns green
- That happens when the last echo or the last explorer arrives




Echo-Algorithm


```
I: {NOT informed} // executed by the initiator
  SEND <Explorer> TO all Neighbors;
  informed := TRUE;

R: {a message from neighbor N is received}
  IF NOT informed THEN // Must be the first explorer
    SEND <Explorer> TO all Neighbors except N;
    informed := TRUE;
    A := N;
  FI
  Count := Count + 1;
  IF Count == #Neighbors THEN
    IF NOT Initiator THEN
      SEND <Echo> TO Neighbor A;
    ELSE
      EXIT; // Algorithm has terminated
    FI
  FI
FI
```


Initially, informed == false
and Count == 0 for all nodes




Echo-Algorithm – Example



Here, the number of already received explorers *and* echos is counted.



Echo-Algorithm – Example




Echo-Algorithm – Characteristics

Exactly two messages run over every edge

- Either an explorer and an echo running in the opposite direction or two explorers running in opposite directions

Parallel traversing of a (connected non-directional) graph with $2e$ messages

- Every node sends an explorer on all edges $\rightarrow +2e$ explorer
- Exception activation edge $\rightarrow -n$ explorer
- Exception initiator $\rightarrow +1$ explorer
- Every node sends an echo on the activation edge $\rightarrow +n$ echos
- Exception initiator $\rightarrow -1$ echo



Echo-Algorithm – Characteristics


The Echo-algorithm is a wave algorithm

Forth wave: becoming red

- Distribution of information (to all nodes over all edges)


Back wave: becoming green

- Collecting of information
(of potentially all nodes over the activation edges)




Echo-Algorithm – Characteristics

- Echo-edges form a spanning tree
- Depending on the message delays, the spanning tree looks differently because fast edges are preferred




Improvement of the Echo-Algorithm?


- Idea: Avoid the visit of nodes which are known to be visited by other explorers
- Together with an explorer, a set of taboo nodes z is sent and received
- The sent taboo set by the initiator is
 $z = \langle \text{neighbors of initiator} \rangle \cup \langle \text{initiator} \rangle$
- Explorers only sent to the set of neighbors y which are not in z .
- Thus, the new taboo set $z' = z \cup y$ is attached
- Advantage: Saving of messages
 - Extreme cases: tree and complete graph
- Disadvantages:
 - message length $O(n)$
 - identity of neighbors has to be known



E.g. the message of 2 and 3 contains the info that 4 does not have to be visited.




Broadcast




Broadcast on Special Topologies

- Broadcast: Sending of a message to all nodes, optionally also with confirmation
- Flooding realizes a broadcast on arbitrary connected undirected topologies
 - Especially fault-tolerant because all edges are used for the distribution of information
- For special topologies, a broadcast with less messages is possible, provided the algorithm knows which topology is underlying
 - Less error-tolerant because, in the aimed case, each node is only reached over a single edge $\rightarrow n - 1$ messages
- Exemplary topologies: Rings, trees, hypercubes



Broadcast on Unidirectional Rings

- Token circulates with message
- All nodes are informed, if the token reaches the initiator again
- n messages
- A ring can also be overlaid by another topology
→ logical ring



Broadcast on Trees


Tree has $n - 1$ edges

One message goes over each edge

For the confirmation (if required) one additional message goes over every edge

Tree can be overlaid by another topology

→ Spanning tree



Broadcast on Hypercubes?

Initiator may have number 00...00 (binary)






Broadcast on Hypercubes

Initiator may have number 00...00 (binary)

Analogous to recursive construction of a hypercube

- Initiator sends in dimension 1
- Then all nodes of dimension 1 in dimension 2
- Then all nodes of dimension 2 in dimension 3
- ...




Broadcast on Hypercubes

Unit Time Complexity

- After d cycles all nodes are informed
- Is that optimal?

Message complexity

- $1 + 2 + 4 + \dots + 2^{d-1} = 2^d - 1 = n - 1$
- Is that optimal?



Literature

E. Chang. *Echo algorithms: Depth parallel operations on graphs*. IEEE Transactions on Software Engineering, 8(4):391--400, 1982.

