

# Distributed Algorithms 2015/16

## **Fault Tolerance**

Reinhardt Karnapke | Communication and Operating Systems Group

---


# Overview

- Introduction into fault tolerance (this lecture)
- Masking fault tolerance (next lecture)
  - Redundancy techniques (only sketched, c.f. lecture “Dependable Systems”)
  - Consensus and related problems
- Non-masking fault tolerance (the next but one lecture)
  - Self-Stabilization



# Fault Tolerance

- No (non-trivial) system contains no fault!
- Taking faults into account is, thus, absolutely necessary!
- In large systems (e.g., the Internet) some components will always work faulty
- Simple motivation: All computers of a systems must be available at the same time (→ **serial composition**)
  - 1 computer → system unavailable 1% of the time
  - 10 computers → system unavailable ~10% of the time
  - 100 computers → system unavailable ~63% of the time
  - 1000 computers → system unavailable ~99.99% of the time
- Example university computer room: How often are **all** computers in a large room functional at the **same time**?



# Basic Terms

## Fault


- Triggering event, e.g., caused by external disturbance or wear/abrasion

## Error

- Internal system state that does not fit the specification, caused by a fault


## Failure

- System does not provide the correct service to the outside, caused by an error
- This chain might abort after every step, a failure is not always the consequence




## Classification of Faults– Benign Faults

- **Crash Fault:** A node suddenly fails and afterwards no actions are exercised, but until the moment of the fault it behaves according to its specification,
  - Also denoted **Halting fault** or **Fail Stop** etc. Often, those terms additionally mean that the fault can be surely determined by all nodes that have not failed
  - **Omission Fault:** Some actions are not exercised
  - **Timing Fault:** Specification suitable behavior, but some actions are carried out too late
  - All have in common: The faulty process does not exercise any actions that a correct process would not do
- ⇒ Not always realistic!

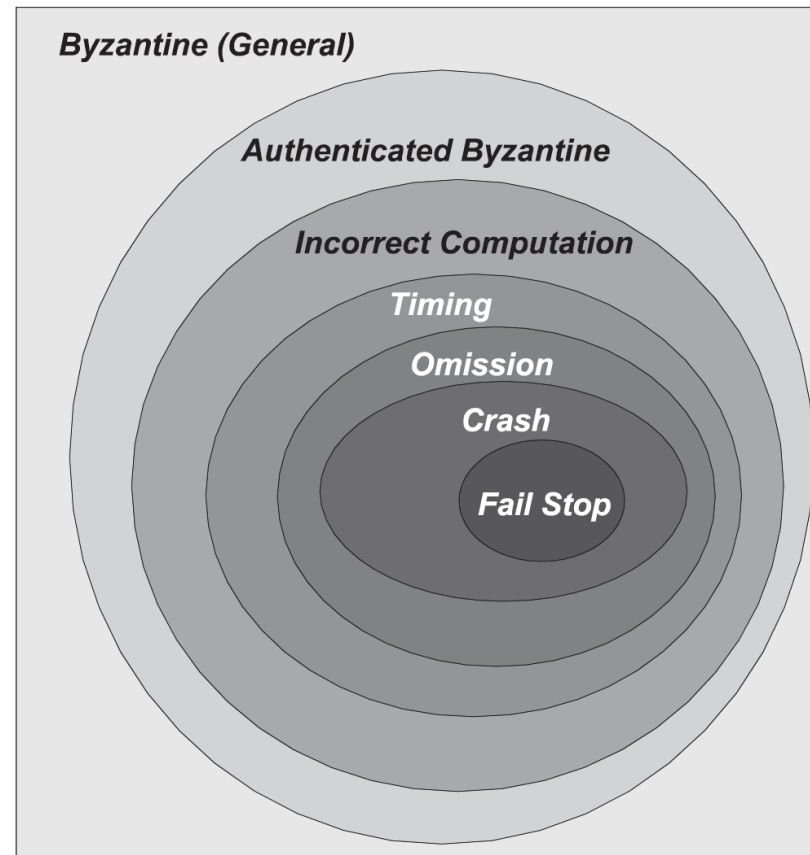


## Classification of Faults– Malicious Faults

- Often denoted as **Byzantine faults** (Lamport, 1982)
- Faulty processes can exercise arbitrary actions and cooperate among each other
  - For example, exercise arbitrary calculations and send arbitrary messages
- Often, also limited model
  - E.g., faulty processes can exercise arbitrary actions calculable with polynomial complexity (i.e., they cannot fake digital signatures of correct nodes)
- Model covers all kinds of faults, e.g., it also targets attacks on a system from outside
- Model also contains all “simpler” cases




## Example: Fault models for Distributed Systems



# Approaches


- Fault intolerance
  - Idea: “Avoid faults!”
  - Eliminating fault causes
    - Very reliable components
    - Extensive testing
    - ...
- Fault tolerance
  - Idea: “Faults occur, we tolerate them!”
  - Fault tolerance by redundancy in space and/or time
  - Acceptance of partial or temporal failures
  - → Considered here





# Paradigms of Fault Tolerance

- **Masking fault tolerance**
  - Aim: avoid system failure if possible
- **Non-masking fault tolerance**
  - System may fail partly or temporarily
  - Better than a complete and/or permanent failure




# Masking Fault Tolerance

- Necessary if a (also temporary) failure of the system would have unacceptable consequences
  - e.g., death of humans or high financial losses
- Sensible also with many other applications
- Tries to ensure safety and liveness!
- Example car brakes
  - Separated brake circuits for right front wheel and left back wheel as well as for left front wheel and right back wheel
  - Car can still break if one circuit has failed



# Masking Fault Tolerance

- Always needs **redundancy** for implementation
- Always only possible for the faults considered
- Can never take into account all possible faults
- Only successful if only a limited part of the components is erroneous
- The amount of the erroneous components that can be tolerated depends on the **fault model**



# Redundancy in Space or Time

- **Redundancy in Space:**


multiple instances of components

- example: a server has several independent power supplies
  - single, functioning power supply sufficient
  - erroneous power supplies are substituted without disrupting server operation

- **Redundancy in Time:**

multiple execution of actions


- example: packages are submitted several times over an unreliable network; from those packages, an error free package is generated (if possible)



# Active and Passive Replication


- **Active Replication**

- All replicas collaborate productively
  - e.g. eyes, aircraft engines, brake systems,...
- TMR (*Triple Modular Redundancy*)



- **Passive Replication**


- Replicas only get active in case of an error (*primary/backup*)
- e.g., spare tire (no masking), emergency power generator



# Fault Tolerance by Redundant Components


- ***k*-reliability** of a system with respect to a component that exists  $n$  times: system can tolerate the failure of up to  $k$  of the  $n$  instances of the components
- Non-Byzantine fault
  - System with an  $n$ -times existing component is  $(n - 1)$ -reliable in reference to that component
- Byzantine fault
  - Assumption: correct output can uniquely be determined from the  $n$  outputs by highly reliable voter
  - System with an  $n$ -times existing component is

$$\left\lfloor \frac{n - 1}{2} \right\rfloor \text{ reliable in reference to that component}$$




## Example: 1–Reliability

- Consider a file server *S* with the operations *read()* and *write()*
- Non-byzantine fault
  - *S* sends, in response to *read()*, either correct content or nothing
  - Realization with two servers *S1* and *S2*
    - *write()*: to both servers
    - *read()*: if *S1* does not reply, then *S2*  
also possible: ask *S1* and *S2*; use first answer
- Byzantine fault
  - *S* sends either the correct content or a false content or nothing after *read()* is called
  - Realization with three servers *S1*, *S2* and *S3*
    - *write()*: to all three servers
    - *read()*: to all three servers, majority vote at the client



## Example: *N*-Version Programming


- Application is implemented several times by *independent* programmer teams
- Results are evaluated by majority vote at run time
- Problems
  - Comparison of results with admissible inaccuracy (e.g., result of a numeric approximation)
  - Comparison of multiple possible correct solutions (e.g., zero of a polynomial)
  - Development often not really “independent”
  - ...





## Problem: Hidden Dependencies

- Hidden dependencies of redundant instances increase probability of simultaneous failure!
- Examples?
  - Redundant power supplies in the same electric circuit
  - Redundant servers cooled by the same air condition
  - Redundant diesel generators sharing the diesel tank
  - Programming teams with same education in N-version programming
  - Same, faulty compiler in N-version programming
  - Redundant computation centers in the same area
  - ...



# Non-Masking Fault Tolerance

- Applicable if a partial or temporary failure of the system/the component is acceptable
- 1. possible aim: In case of an error, bring the system into a safe state (**fail safe**)
  - Assures safety
  - Example: traffic light control
    - In case of an error of the traffic light control, all traffic lights are switched to red
  - Example: mechanic stop sign for trains
    - When the signal rope is torn, the arm of the signal falls into the position “Stop!”




# Non-Masking Fault Tolerance

2. possible aim: in case of a failure, the system is run on with restricted functionality


**(graceful degradation)**

- Assures safety and limited liveness
- Example servo steering
  - In case of the failure of the servo pump, steering is still possible, but only with higher effort
- Example ABS brake
  - In case of electronic failure, the brake is still operational without ABS functionality



## Non-Masking Fault Tolerance

3. possible aim: In case of a failure, reconfigure the system in such a way that it works correctly again
- Assures safety again after reconfiguration (**eventual safety**)
  - Liveness is not affected if the reconfiguration only lasts for a limited time
  - Example: communication over serially connected lines
    - In case of the failure of one or several lines, an alternative route is set up without usage of the failed lines



## Literature

1. T. Anderson, P.A. Lee: Fault Tolerance – Principles and Practice, Prentice Hall, 1982
2. D.K. Pradhan (Hrsg.): Fault Tolerant Computer Systems, Prentice Hall, 1996
3. D.P. Siewiorek, R.S. Swarz: The Theory and Practice of Reliable Systems Design, Digital Press, 1995

