

## Model-based Collaborative Filtering

Neighborhood-based approaches for collaborative filtering are specific to the user and item that is being predicted. In contrast, **model-based collaborative filtering** relies on a summarized model of the data that is created beforehand during a training phase. This model is then used in the later prediction phase to efficiently calculate the predictions.


Advantages of model-based approaches<sup>1</sup>:

- ▶ Space-efficiency: Typically, the size of the learned model is much smaller than the original ratings matrix.
- ▶ Training speed and prediction speed: One problem with neighborhood-based methods is that the pre-processing stage is quadratic in either the number of users or the number of items. Model-based systems are usually much faster in the preprocessing phase of constructing the trained model.
- ▶ Avoiding overfitting: Overfitting is a serious problem in many machine learning algorithms, in which the prediction is overly influenced by random artifacts in the data. The summarization approach of model-based methods can often help in avoiding overfitting.

---

<sup>1</sup>Adopted from: Recommender Systems The Textbook, C. C. Aggarwal, Springer, 2016

# Model-based Collaborative Filtering



Comparing the traditional classification problem with collaborative filtering. Shaded entries are missing and need to be predicted<sup>1</sup>

Model-based collaborative filtering methods include, for example, rule-based methods, decision trees, regression models, Bayes classifiers, support vector machines, or neural networks.

<sup>1</sup>Source: Recommender Systems The Textbook, C. C. Aggarwal, Springer, 2016

## Association Rule Mining

Association Rule Mining is a technique to identify rulelike relationship patterns in large-scale transactions.

- ▶ An example rule could be: "If a customer purchases baby food then he or she also buys diapers in 70 percent of the cases".

Definition by Sarwar et al.:


- ▶ A (sales) transaction  $T$  is a subset of the set of available products  $P = \{p_1, \dots, p_m\}$  and describes a set of products that have been purchased together.
- ▶ Association rules are often written in the form  $X \rightarrow Y$ , with  $X$  and  $Y$  being both subsets of  $P$  and  $X \cap Y = \emptyset$
- ▶ An association rule  $X \rightarrow Y$  (e.g. baby food  $\rightarrow$  diapers) expresses that whenever the elements of  $X$  (the rule body) are contained in transaction  $T$ , it is very likely that the elements in  $Y$  (the rule head) are elements of the same transaction.

In collaborative recommender systems, the goal of association rule mining is to detect rules such as *"If user Alice liked both Item 5 and Item 7, then Alice will most probably also like Item 23"*.

## Association Rule Mining


The goal of rule mining algorithms such as *Apriori* (by Agrawal and Srikant, 1994) is to automatically detect such rules and *calculate a measure of quality for those rules*<sup>1</sup>.

- The **support** of a rule  $X \rightarrow Y$  is calculated as the percentage of transactions that contain all items of  $X \cup Y$  with respect to the number of overall transactions (i.e., the probability of co-occurrence of  $X$  and  $Y$  in a transaction):



$$\text{Support}(X \rightarrow Y) = \frac{\text{Number of transactions containing } X \cup Y}{\text{Number of transactions}}$$

- The **confidence** is defined as the ratio of transactions that contain all items of  $X \cup Y$  to the number of transactions that contain only  $X$ :



$$\text{Confidence}(X \rightarrow Y) = \frac{\text{Number of transactions containing } X \cup Y}{\text{Number of transactions containing } X}$$

<sup>1</sup>Adapted from: D. Jannach et al.: Recommender Systems An Introduction, Cambridge University Press, 2011

## Association Rule Mining: Example

Assume a ratings matrix with a binary scale ("like" / "dislike"):

	Bread	Butter	Milk	Fish	Beef	Ham
User 1	1	1	1	0	0	0
User 2	0	1	1	0	1	0
User 3	1	1	0	0	0	0
User 4	1	1	1	1	1	1
User 5	0	0	0	1	0	1
User 6	0	0	0	1	1	1
User 7	0	1	0	1	1	0

It is evident that the columns of the table can be partitioned into two sets of closely related items. One of these sets is  $\{Bread, Butter, Milk\}$ , and the other set is  $\{Fish, Beef, Ham\}$ . These are the only itemsets with at least 3 items, which also have a support of at least 0.2.

## Association Rule Mining

Finding association rules involves two steps:

1. Find all itemsets that satisfy a minimum support threshold  $s$ .
2. From each of these itemsets  $Z$ , all possible 2-way partitions  $(X, Z - X)$  are used to create a potential rule  $X \rightarrow Z - X$ . Those rules satisfying the minimum confidence are retained.

The calculation of interesting association rules can be performed offline.

- ▶ At runtime, the following scheme can be used to compute recommendations for a user Alice:
  - ▶ Determine the set of  $X \rightarrow Y$  association rules that are relevant for Alice that is, where Alice has bought (or liked) all elements from  $X$ .
  - ▶ Compute the union of items appearing in the consequent  $Y$  of those association rules that have not been purchased by Alice.
  - ▶ Sort the products according to the confidence of the rule that predicted them.
  - ▶ Return the first  $N$  elements of this ordered list as a recommendation.



## Matrix factorization / Singular Value Decomposition

- ▶ Simply put, matrix factorization methods can be used in recommender systems to derive a set of **latent** (hidden) factors from the rating patterns and characterize both users and items by such vectors of factors.
- ▶ **Singular Value Decomposition (SVD)** was proposed by Deerwester et al. in 1990 as a method to discover latent factors in documents.
- ▶ The SVD theorem by Golub and Kahan (1965) states that a given matrix **M** can be decomposed into a product of three matrices as follows:

$$\mathbf{M} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$$



where **U** and **V** are called *left* and *right singular vectors* and the values of the diagonal of  $\mathbf{\Sigma}$  are called the *singular values*.

## Singular Value Decomposition: Example

Consider the following rating matrix:

	Item 1	Item 2	Item 3	Item 4
User 1	3	1	2	3
User 2	4	3	4	3
User 3	3	2	1	5
User 4	1	6	5	2

Because in this example the  $4 \times 4$  matrix  $\mathbf{M}$  is quadratic,  $\mathbf{U}$ ,  $\Sigma$  and  $\mathbf{V}$  are also quadratic.

$$\mathbf{M} = \mathbf{U}\Sigma\mathbf{V}^T$$

$$\mathbf{U} = \begin{pmatrix} -0.35 & -0.36 & 0.29 & -0.80 \\ -0.56 & -0.08 & 0.62 & 0.52 \\ -0.44 & -0.56 & -0.65 & 0.21 \\ -0.59 & -0.73 & -0.28 & -0.17 \end{pmatrix}$$

$$\Sigma = \begin{pmatrix} 12.22 & 0 & 0 & 0 \\ 0 & 4.92 & 0 & 0 \\ 0 & 0 & 2.06 & 0 \\ 0 & 0 & 0 & 0.29 \end{pmatrix}$$

$$\mathbf{V}^T = \begin{pmatrix} -0.43 & -0.53 & -0.52 & -0.50 \\ -0.49 & 0.53 & 0.40 & -0.55 \\ 0.55 & -0.41 & 0.48 & -0.53 \\ 0.51 & 0.50 & -0.56 & -0.38 \end{pmatrix}$$



## Singular Value Decomposition: Example

The main point of the SVD is that we can approximate the full matrix by observing only the most important features, i.e., those with the largest singular values.

- For example, consider the projection of  $\mathbf{U}$ ,  $\mathbf{V}^T$  and  $\Sigma$  in the two-dimensional space:

$$\mathbf{U}_2 = \begin{pmatrix} -0.35 & -0.36 \\ -0.56 & -0.08 \\ -0.44 & -0.56 \\ -0.59 & -0.73 \end{pmatrix} \quad \Sigma_2 = \begin{pmatrix} 12.22 & 0 \\ 0 & 4.92 \end{pmatrix} \quad \mathbf{V}_2^T = \begin{pmatrix} -0.43 & -0.53 & -0.52 & -0.50 \\ -0.49 & 0.53 & 0.40 & -0.55 \end{pmatrix}$$


- Matrix  $\mathbf{U}$  and  $\mathbf{V}$  correspond to the **latent user and item factors**.
- Although in this example we cannot observe any clusters of users, we can see that the items from  $\mathbf{V}$  build two groups.



## Singular Value Decomposition: Example

Consider a target user *Alice* with a rating vector  $r_{\text{Alice}} = [5, 3, 4, 4]$  for the four items. We can find out where Alice would be positioned in this two-dimensional space by calculating

$$r_{\text{Alice}2D} = r_{\text{Alice}} \times \mathbf{V}_2 \times \Sigma_2^{-1} \approx [-0.63, 0.29]$$



## Naive Bayes Collaborative Filtering

Naive Bayes is a generative model, which is commonly used for classification.

- ▶ We treat **items as features and users as instances.**
- ▶ The main challenge is that any feature (item) can be the target class, and we have to deal with incomplete feature variables.

Assume a small number of  $l$  distinct ratings  $v_1, \dots, v_l$  such as *like*, *neutral*, *dislike*, and an  $m \times n$  rating matrix where the  $(u, j)$ th element is denoted by  $r_{uj}$ . Furthermore, let  $I_u$  be the set of items that have been rated by user  $u$ .

The goal of the Bayes classifier is to predict the unobserved rating  $r_{uj}$ .

- ▶  $r_{uj}$  can take any one of the discrete possibilities in  $\{v_1, \dots, v_n\}$ .
- ▶ We would like to determine the **probability** that  $r_{uj}$  takes on any of these values **conditional on the observed ratings in  $I_u$ .**

In other words, we want to determine the **probability**  $P(r_{uj} = v_s | \text{Observed ratings in } I_u)$  for each value of  $s$  in  $\{1, \dots, l\}$ .

## Naive Bayes Collaborative Filtering

We can simplify this expression by using the well-known **Bayes theorem**:

$$P(A|B) = \frac{P(B|A) * P(A)}{P(B)}$$



Therefore for each value of  $s$  in  $\{1, \dots, I\}$  we have

$$P(r_{uj} = v_s | \text{Observed ratings in } I_u) = \frac{P(r_{uj} = v_s) * P(\text{Observed ratings in } I_u | r_{uj} = v_s)}{P(\text{Observed ratings in } I_u)}$$

- ▶ We would like to determine the value of  $s$  for which the value of  $P(r_{uj} = v_s | \text{Observed ratings in } I_u)$  is as large as possible.
- ▶ Because the denominator of the right hand side is independent of  $s$ , we can express the equation in terms of a constant of proportionality:

$$P(r_{uj} = v_s | \text{Observed ratings in } I_u) \propto P(r_{uj} = v_s) * P(\text{Observed ratings in } I_u | r_{uj} = v_s)$$

## Naive Bayes Collaborative Filtering

- ▶ The value of  $P(r_{uj} = v_s)$ , also called the **prior probability**, is estimated to the fraction of users that have specified the rating  $v_s$  for item  $j$  (users who haven't rated  $j$  are ignored here).
- ▶  $P(\text{Observed ratings in } I_u | r_{uj} = v_s)$  is estimated with the use of the **naive assumption** (i.e., conditional independence between ratings). Conditional independence says that the ratings of a user for various items  $I_u$  are independent from each other, *conditional* of the fact that the value of  $r_{uj}$  was observed to be  $v_s$ . We can express this mathematically as follows:

$$P(\text{Observed ratings in } I_u | r_{uj} = v_s) = \prod_{k \in I_u} P(r_{uk} | r_{uj} = v_s)$$

- ▶  $P(r_{uk} | r_{uj} = v_s)$  is estimated as the fraction of users that have specified the rating of  $r_{uk}$  for the  $k$ th item, *given that they have specified the rating of their  $j$ th item to  $v_s$ .*

## Naive Bayes Collaborative Filtering

By plugging in the estimation of the prior probability  $P(r_{uj} = v_s)$  and the previous equation, it is possible to obtain an estimate of the **posterior probability** of the rating of item  $j$  for user  $u$  as follows:

$$P(r_{uj} = v_s | \text{Observed ratings in } I_u) \propto P(r_{uj} = v_s) * \prod_{k \in I_u} P(r_{uk} | r_{uj} = v_s)$$

By computing each of the expressions on the right-hand side for each  $s \in \{1, \dots, I\}$ , and determining the value of  $s$  at which it is the largest, we can determine the most likely value  $\hat{r}_{uj}$  of the missing rating  $r_{uj}$ :



$$\hat{r}_{uj} = \operatorname{argmax}_{v_s} P(r_{uj} = v_s) * \prod_{k \in I_u} P(r_{uk} | r_{uj} = v_s)$$