

Sheet 2: Maximum Likelihood Estimation

In this exercise sheet, we will look at various properties of maximum-likelihood estimation, and how to find maximum-likelihood parameters.

ML vs. James Stein Estimator (15 P)

Let $X_1, \dots, X_n \in \mathbb{R}^d$ be independent draws from a multivariate Gaussian distribution with mean vector μ and covariance matrix $\Sigma = \sigma^2 I$. It can be shown that the maximum-likelihood estimator of the mean parameter μ is the empirical mean given by:

$$\hat{\mu}_{\text{ML}} = \frac{1}{N} \sum_{i=1}^N X_i$$

It was once believed that the maximum-likelihood estimator was the most accurate possible (i.e. the one with the smallest Euclidean distance from the true mean). However, it was later demonstrated that the following estimator

$$\hat{\mu}_{\text{JS}} = \left(1 - \frac{(d-2) \cdot \sigma^2}{n \cdot \|\hat{\mu}_{\text{ML}}\|^2}\right) \hat{\mu}_{\text{ML}}$$

(a shrunk version of the maximum-likelihood estimator towards the origin) has actually a smaller distance from the true mean when $d \geq 3$. This however assumes knowledge of the variance of the distribution for which the mean is estimated. This estimator is called the James-Stein estimator. While the proof is a bit involved, this fact can be easily demonstrated empirically through simulation. This is the object of this exercise.

The code below draws ten 50-dimensional points from a normal distribution with mean vector $\mu = (1, \dots, 1)$ and covariance $\Sigma = I$.

```
In [1]: def getdata(seed):

    n = 10                # data points
    d = 50                # dimensionality of data
    m = numpy.ones([d])   # true mean
    s = 1.0               # true standard deviation

    rstate = numpy.random.mtrand.RandomState(seed)
    X = rstate.normal(0,1,[n,d])*s+m

    return X,m,s
```

The following function computes the maximum likelihood estimator from a sample of the data assumed to be generated by a Gaussian distribution:

```
In [2]: def ML(X):
    return X.mean(axis=0)
```

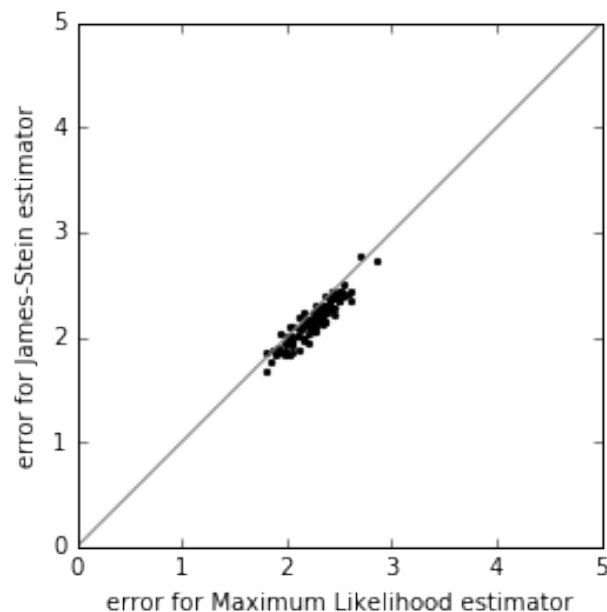
- Based on the ML estimator function, write a function that receives as input the data $(X_i)_{i=1}^n$ and the (known) variance σ^2 of the generating distribution, and computes the James-Stein estimator

```
In [3]: def JS(X,s):
    # REPLACE BY YOUR CODE
    import solution
    m_JS = solution.JS(X,s)
    ###
    return m_JS
```

We would like to compute the error of the maximum likelihood estimator and the James-Stein estimator for 100 different samples (where each sample consists of 10 draws generated by the function `getdata` with a different random seed). Here, for reproducibility, we use seeds from 0 to 99. The error should be measured as the Euclidean distance between the true mean vector and the estimated mean vector.

- Compute the maximum-likelihood and James-Stein estimations.
- Measure the error of these estimations.
- Build a scatter plot comparing these errors for different samples.

```
In [4]: %matplotlib inline
        ### REPLACE BY YOUR CODE
        import solution
        solution.compare_ML_JS()
        ###
```



Parameters of a mixture of exponentials (15 P)

We consider the following “mixture of exponentials” distribution supported on \mathbb{R}^+ , that we use to generate data, but whose parameters α and β are unknown.

$$p(x; \alpha, \beta) = 0.5 \cdot [\alpha e^{-\alpha x} + \beta e^{-\beta x}]$$

A dataset $\mathcal{D} = x_1, \dots, x_N$ with $N = 200$ has been generated from that distribution. It is given below and plotted as a histogram.

```
In [5]: D=[ 0.74,  0.20,  0.56,  0.05,  0.67,  0.41,  0.74,  4.63,  0.59,  0.39,
            0.71,  0.17,  5.34,  0.33,  0.01,  1.11,  0.60,  0.41,  0.65,  1.97,
            0.19,  0.80,  0.04,  0.48,  0.54,  0.59,  0.31,  1.40,  0.63,  0.38,
            0.36,  0.02,  0.68,  0.72,  0.84,  0.30,  0.01,  1.37,  0.89,  0.10,
            0.21,  0.68,  0.14,  0.10,  0.11,  0.01,  0.09,  0.50,  0.34,  0.30,
            1.22, 10.05,  0.19,  0.04,  0.13,  1.53,  2.28,  1.76,  0.03,  0.31,
```

```


0.37, 0.50, 0.05, 0.30, 0.53, 0.63, 4.20, 0.86, 0.29, 1.98,
1.27, 0.35, 0.43, 0.35, 0.75, 0.25, 1.15, 1.65, 0.82, 0.37,
2.55, 2.75, 3.06, 0.97, 2.65, 8.97, 0.04, 2.98, 0.36, 0.01,
0.85, 0.90, 0.09, 0.01, 0.82, 2.30, 2.09, 0.29, 0.16, 2.12,
5.28, 0.27, 0.15, 1.02, 0.51, 0.02, 1.72, 1.35, 0.51, 0.27,
1.05, 2.24, 3.93, 0.62, 3.38, 0.56, 0.49, 2.84, 0.27, 0.12,
3.99, 0.16, 0.09, 3.61, 0.54, 0.08, 0.31, 1.38, 0.63, 0.61,
0.21, 0.13, 2.28, 2.61, 4.60, 0.02, 0.34, 0.15, 0.07, 2.44,
0.86, 0.73, 2.01, 0.26, 0.72, 1.56, 0.09, 0.97, 0.24, 0.92,
1.05, 0.71, 1.28, 3.79, 1.32, 0.17, 0.39, 2.82, 0.12, 2.06,
2.04, 0.00, 1.94, 0.27, 0.91, 0.36, 0.92, 5.69, 0.33, 0.69,
1.00, 2.19, 0.01, 0.08, 1.16, 0.31, 0.83, 0.41, 1.27, 0.08,
4.69, 0.65, 0.43, 0.10, 2.92, 0.06, 6.21, 0.90, 0.00, 0.52,
0.65, 0.26, 1.94, 0.37, 0.50, 5.66, 4.24, 0.40, 0.39, 7.89]

```

```

%matplotlib inline
from matplotlib import pyplot as plt
plt.hist(D,bins=30)
plt.show()

```



For this dataset, the log-likelihood function is given by


$$\ell(\alpha, \beta) = \log \prod_{i=1}^N p(x_i; \alpha, \beta) = \sum_{i=1}^N \log(e^{-\alpha x_i} + \beta e^{-\beta x_i}) - \log(2)$$

Unfortunately, it is difficult to extract the parameters α, β analytically by solving directly the equation $\nabla \ell = 0$. Instead, we will analyze the function over a grid of parameters α, β . We know a priori that parameters α and β are in the intervals $[0.4, 1.0]$ and $[1.5, 4.5]$ respectively.

- **Build a grid on this limited domain and evaluate log-likelihood at each point of the grid.**
- **Plot the log-likelihood function as a contour plot, and superpose the grid to it.**

Highest log-likelihood values (i.e. most probable parameters) should appear in red, and lowest values should be plotted in blue. Two adjacent lines of the contour plot should represent a log-likelihood difference of 1.0. In your code, favor numpy array operations over Python loops.

```
In [6]: ### REPLACE BY YOUR CODE
import solution
solution.s2a(D)
###
```



Gradient-Based Optimization (10 P)

As an alternative to computing the log-likelihood for a whole grid, we would like to find the optimal parameters α, β by gradient-based optimization. The partial derivatives of the log-likelihood function are given by:

$$\frac{\partial \ell(\alpha, \beta)}{\partial \alpha} = \sum_{i=1}^N \frac{e^{-\alpha x_i} (1 - \alpha x_i)}{\alpha e^{-\alpha x_i} + \beta e^{-\beta x_i}}$$

$$\frac{\partial \ell(\alpha, \beta)}{\partial \beta} = \sum_{i=1}^N \frac{e^{-\beta x_i} (1 - \beta x_i)}{\alpha e^{-\alpha x_i} + \beta e^{-\beta x_i}}$$


A gradient ascent step of the log-likelihood function takes the form

$$\begin{pmatrix} \alpha \\ \beta \end{pmatrix} \leftarrow \begin{pmatrix} \alpha \\ \beta \end{pmatrix} + \gamma \nabla_{\alpha, \beta} \ell(\alpha, \beta)$$

where γ is a learning rate to be defined. We start with initial parameters $\alpha = 0.7$ and $\beta = 3.0$.

- Implement the gradient ascent procedure.
- Run the gradient ascent with parameter $\gamma = 0.005$.
- Plot the trajectory of the gradient ascent in superposition to the contour plot of the previous exercise.

```
In [7]: ### REPLACE BY YOUR CODE
import solution
solution.s2b(D)
###
```



As it can be seen, the optimization procedure does not converge in reasonable time and seems to oscillate.

- **Explain the problem(s) with this approach. Propose a simple improvement of the optimization technique and apply it.**

[REPLACE BY YOUR EXPLANATION + PROPOSITION]

```
In [8]: ### REPLACE BY YOUR CODE
import solution
solution.s2c(D)
###
```

